

ARC BUILDERS FUND

Machines are starting to earn. We built the part that **proves they did the work.**

StonkRobotics — the evaluation and settlement layer for
machine work on Arc.

Richard, Founder · 2026-09-20 · stonkrobotics.xyz · github.com/robot-policy-network · Builders
Fund consideration

THE PROBLEM

Every existing answer fails in a way that matters.

APPROACH	WHAT ACTUALLY HAPPENS
Agent payment tools	A human still decides. Nothing is verified — the wallet is automated, not the trust.
Oracle / escrow designs	The reported score <i>is</i> the payment instruction, so a manipulated or hallucinated result moves money.
Human review	Works, does not scale, and sits slower than the work it reviews.

The missing primitive: **a payment triggered by evidence of work, not by a claim of work — at machine cadence.**

THE DESIGN

```
evaluator signs EvaluationAttestation → evidence  .  
payer      signs ReleaseAuthorization → authority  .  
  
release() executes only when BOTH validate
```

The evaluator key is a **hot key on our server**. The design assumes it will be compromised and makes that specific failure survivable: a compromised evaluator can forge a result on an open job, but **still cannot pay itself** — the payer's separate signature is required.

What this does *not* claim: that the score is trustless. See the next two pages.

SETTLEMENT — LIVE, AND HONESTLY SCOPED

STEP	TRANSACTION
Deploy <code>EvaluationEscrow</code>	<code>0x479FF86C25d813cD3FA076e2Fe6df2E3d2d77491</code>
<code>createJob</code> — 0.5 USDC escrowed	<code>0x087d419b...</code>
<code>verifyResult</code> — evaluator attestation, no funds move	<code>0xad894353...</code>
<code>release</code> — payer authorization, USDC moves	<code>0x4a8ccd2d...</code>
Result	provider 0 → 500000 · job state 3 (RELEASED)

What this proves: the full path executes with real USDC.

What it does not prove: in this demo the payer, evaluator and provider are **all founder-controlled keys**, and the evaluator is the *same address* as the provider — so it demonstrates the process, not the security property against a hostile counterparty.


No independent audit. This instance is a demo; a production deployment uses a fresh owner key (M6).

TRUST BOUNDARY AND KNOWN LIMITS

LIMITATION	STATUS
The evaluator is a single centralized server. "Reproducible" means the scorer is deterministic — not that a third party can recompute it today (simulator unreleased, reference server-side)	recomputation is M5
p1anHash is not recomputed on-chain. A payer that skips off-chain recomputation trusts the evaluator bound the right plan	accepted, spec §10
A payer who never signs leaves funds escrowed. The provider has no on-chain claim and no recourse path	deliberate v2 trade
Near-copies can win. Only an <i>exact</i> copy is neutralised; a copy plus a marginal tweak can be strictly better	known, bounded by unpublished plans
Weak pricing signal inside the passing band. A generic plan lands 95–97; the pass mean is 95.8	known

The accurate summary of the security claim is "**an evaluator compromise cannot move funds**" — not "trustless work". We publish this list because a reviewer will find it anyway.

TRACTION — THE FIVE DAYS



Sep 20 is a partial day. 65,138 scored rounds · 217 wallets · 67 real robot tasks · 80 of 120 hourly windows had activity.

TRACTION — WHAT THE SCORE DOES, AND DOES NOT, SEPARATE

83–87

the hardest real-trajectory task, where 184 distinct wallets landed — evidence of gradation *across* missions

95.8 ± 2.9

the passing mean and spread — a generic plan lands 95–97, so the score barely separates a template from a competent agent *within* the band

99.94% pass rate is **by race design** (agents retry until they clear the threshold), so pass/fail carries little information. We report the gradation and the weakness in the same breath rather than claiming a sharp discriminator.

TRACTION — THE COMMUNITY, INCLUDING THE UGLY PART

SIGNAL	VALUE
Wallets active all 5 days	123 of 217 (57%) · 82% returned on 4+ days
Median rounds per address	297 · median answer diversity 0.988
The four heaviest wallets	~1,350 rounds each, ~26-second cadence, answer diversity $\approx$0.09 — cache replay, not planning
New addresses after day one	9 (208 of 217 arrived on launch day)
Third-party USDC settled · revenue	0 · \$0

The 0.988 median diversity is real and so is the 0.09 at the top of the table. Both are in the published dataset; we show them together.

MECHANISM — EVALUATION

$$\text{score} = 100 \times (\text{seedsClean} / n) \times (0.38 \cdot \text{align} + 0.32 \cdot \text{pick})$$

Agents submit a structured action plan. A pure-function simulator executes it against a world state across several perturbation seeds and returns the traced path, the targets actually serviced, and any constraint violations.

Same input → same output, always. Seeded PRNG, no clock, no global randomness. The weights are a tuned heuristic chosen by the founder, not derived from first principles — stated plainly because a reviewer will ask why 0.38 and not 0.40.

MECHANISM — EACH TERM CLOSES AN ATTACK

ATTACK	CLOSED BY
Submit a plan that does not run	<code>seedsClean/n</code> is multiplicative — a hard violation on any seed cuts the whole score; on all seeds it is zero
Draw a straight line between targets	<code>align</code> = coverage + progression . Coverage alone lets a straight line score 0.95
Service the right targets in the wrong order	<code>pick</code> = LCS on the sequence , not membership
Reuse a winning plan on the next mission	<code>precision</code> uses per-target constants (<code>offset=stem_mm</code> , <code>force=2·stem_mm·(2-ripe)</code>) — different targets, different numbers
Copy a plan inside the same window	the race closes on a strictly better score — an exact copy cannot displace its original
Drive extra laps for efficiency marks	<code>eff</code> is an asymmetric ratio capped at 1 — undershooting is free, overshooting costs
Reverse-engineer the rubric	feedback names only public-state terms, never reference coordinates
Arbitrage integer rounding	ranking uses 2-decimal scores; a generic plan lands 95–97, which would otherwise be a mass tie

Known gap, stated: an exact copy is neutralised, a **near-copy plus a marginal tweak** is not. Plans are not published per-window, which limits but does not eliminate it.

MECHANISM — THE COMPETITION

Because first-past-the-post made it a latency contest.

An agent

closes the loop in ~90 seconds.

A person on a phone

pasting a brief into a chat model
needs 2–3 minutes.

So a human could never win one, however good the plan. Deciding at the close makes submission time worth **nothing** and plan quality worth **everything** — the only version where the two compete in one lane.

MECHANISM — THE ADVERSARIAL ECONOMICS

CONTROL

WHY THIS ONE AND NOT THE OBVIOUS ONE

**Per-source (IP)
daily cap: 4**

Wallet addresses are free and the holdings gate reads 0 forever once a winner stops minting. Per-source is the cap that bites.

**Per-window try
cap: 8**

Bounds brute-forcing plan variants — and every attempt still lands in the dataset as a graded round.

**Per-wallet daily
cap: 2**

Deliberately low. A higher cap does not reward a good agent — it **deletes free slots**. Three wallets took 16 of the first 31 windows and turned them into 3 slots.

**Same mission
per window**

Hard-mode operators do not get a private harder draw, or the window's scores would not be comparable.

MARKET

\$305B

stablecoin market cap · USDC
\$74.4B across 157 chains

DefiLlama, live —
defillama.com/stablecoins

\$7.8B → \$52.6B

AI agents market, 2025 → 2030 (46.3%
CAGR)

MarketsandMarkets —
marketsandmarkets.com/PressReleases/ai-agents.asp

The value flows and the agents that earn them are both real and both growing. The trust layer that lets the first pay the second does not exist.

What we have not sized: the addressable market for a settlement layer specifically. We have a wedge — a live arena, 217 wallets, one executed settlement — not a bottom-up SOM. We would rather say that than invent a number.

BUSINESS MODEL — AND WHAT DOES NOT EXIST YET

STREAM	MECHANISM	TRUTH
Settlement fee	A take rate on USDC released	Not possible in the deployed contract — <code>release()</code> transfers the full escrowed amount. A fee path is a v3 change.
Priority verification	Paid tier above the free race window	free tier only; 0 paid entries
Curated evaluation data	Licensing (task, plan, score) triples	dataset published; no buyer

Revenue today: **\$0**. These are architectural possibilities, and the arena is the price-discovery vehicle — not current income.

COMPETITION

PLAYER	RAISES	LAYER	GAP WE FILL
Axis Robotics	\$12M seed (Hack VC)	collection engine	no verification protocol, no settlement
PrismaX	\$11M (a16z CSX)	collection + teleop + models	internal scoring, not on-chain verifiable
RoboTrain (Virtuals)	Virtuals launch	teleop + human scoring	centralized grading, no crypto settlement
StonkRobotics	—	evaluation → verification → settlement	the intersection, on Arc mainnet

Funding figures as reported in company announcements and press. We compete on the layer none of them occupies — and the funded players could enter it next, which is why the on-chain record matters more than the code.

THE MOAT

- **A deterministic, execution-anchored evaluator** — most "scoring" projects use a human panel or a non-deterministic LLM call, which cannot anchor an on-chain artifact
- **An on-chain evidence/authority split** — deployed and executed; a protocol-level property, not a UI choice
- **Arc-native settlement** — low-cost, fast-finality, native USDC

Each alone is copyable; the three together, already frozen and running, are harder. But the honest moat is not the code — it is open-source. It is the **time-stamped record of verified outcomes** and the operator community that a late entrant cannot backfill.

Counterweight: on-chain, that record is currently ~60 NFT mints and **one** founder-funded settlement. The 65,138 rounds live in a published but operator-controlled dataset. The moat is early, not built.

WHY ARC

- **Product reason:** the settlement leg needs a stablecoin that is the gas asset and a fast-finality chain, so a machine payment is one asset and one hop. Arc gives native USDC with an ERC-20 interface over the same balance.
 - **Integration reason:** we hit Arc's USDC compliance precompile (`0x1800...0001::isBlocklisted`) inside `transferFrom` and documented what it breaks in local simulation vs. real execution — field notes other Arc builders can use.
 - **Not a reason:** grant fit. AI2Human is on Base because its users are there. This protocol is on Arc because that is where machine settlement is cheapest.
-

WHAT WE ARE RAISING

WHAT	DETAIL
What we are seeking	Capital from Circle Ventures, deployed at their sole discretion. We are not naming a number — we are open to discussing.
What the capital does	Accelerates M3 (Circle Agent Stack integration) and M4 (independent security review). Both are already scoped, already specified, and the codebase is already on mainnet.
What Circle gets	A production-grade settlement primitive that is already live on Arc mainnet, plus a front-row seat in defining how machines settle payments on Arc.

This is a founder-funded demo today. The capital does not fund an idea — it funds the acceleration of something that already runs.

TEAM

Richard — solo founder.

AI2Human Network — verification layer for open-world AI agents: *task* → *proof* → *verify* → *settle*. Live on Base. GitHub org [ai2humannetwork](https://github.com/ai2humannetwork). **2nd prize, X Layer (OKX) hackathon.**

That project proves the thesis on **human** execution. This one proves it on **machine** execution — and this one is already on mainnet.

Contract deployed · settlement executed · data public and independently checkable.

SCOPE, DISCLOSURES, LINKS

Site · agent workflow	stonkrobotics.xyz · stonkrobotics.xyz/skill/
Repository	github.com/robot-policy-network/robot-policy-network
Settlement proof (every tx)	deployments/arc-mainnet-settlement.json
Dataset + verifier	docs/grants/traction-evidence/verify.sh
Trust limits · milestones · mechanism	docs/grants/MILESTONES-CANONICAL.md · docs/mechanism-deep-dive.md
Not part of this proposal	ARCROBO token and the UnitreeG1Fleet NFT collection are separate subsystems on the same chain — the free-mint competition that produced the traction ran through the NFT collection, not the settlement protocol